

Reinforcement Learning Fundamentals

Lecture 2: RL Framework

Dr Sandeep Manjanna

Assistant Professor, Plaksha University

sandeep.manjanna@plaksha.edu.in

Some material in this lecture is taken from

1. Prof. Ravindran's course: "Reinforcement Learning."
2. Dr Silver's course: "Reinforcement Learning."



Course Logistics


- Classes on **M, W, and F** at **12:00 to 12:50 pm** in Room **1001**.
- Office hours: **Mondays 3:00 to 4:00 pm** (Office **2407**).
- TFs → Keshav Sivakumar Office Hour: Wed 11:00 to 12:00 pm
 Poonam Adhikari Office Hour: Mon 2:00 to 3:00 pm

What is RL?

- Learning about stimuli and actions based on rewards and punishments alone.
- No detailed supervision available
- Trial-and-error learning
- Delayed rewards
- Sequence of actions required to obtain reward
- Associative learning required
 - Need to associate actions to states
- Learn about policies not just actions
- Typically in a stochastic world

Applications of RL

- Stanford's autonomous helicopter flight: extreme aerobatics under computer control (2008)



Source: <http://heli.stanford.edu/>


Applications of RL

- McGill's Adaptive Gait Control for Swimming Robots (2015)

Learning Legged Swimming Gaits from Experience

ICRA 2015 - Best Paper Award Nominee

http://www.cim.mcgill.ca/~dmeger/ICRA2015_GaitLearning/



David Meger, Juan Camilo Gamboa Higuera,
Anqi Xu, Philippe Giguere and Gregory Dudek

Source:

<https://www.youtube.com/watch?v=VHmEdRXSi9M>

Applications of RL

- Gerald Tesauro at IBM's Thomas J. Watson Research Center on TD Gammon (Backgammon game) (1992)



Applications of RL

- Ref: Playing Atari with Deep Reinforcement Learning (2013)


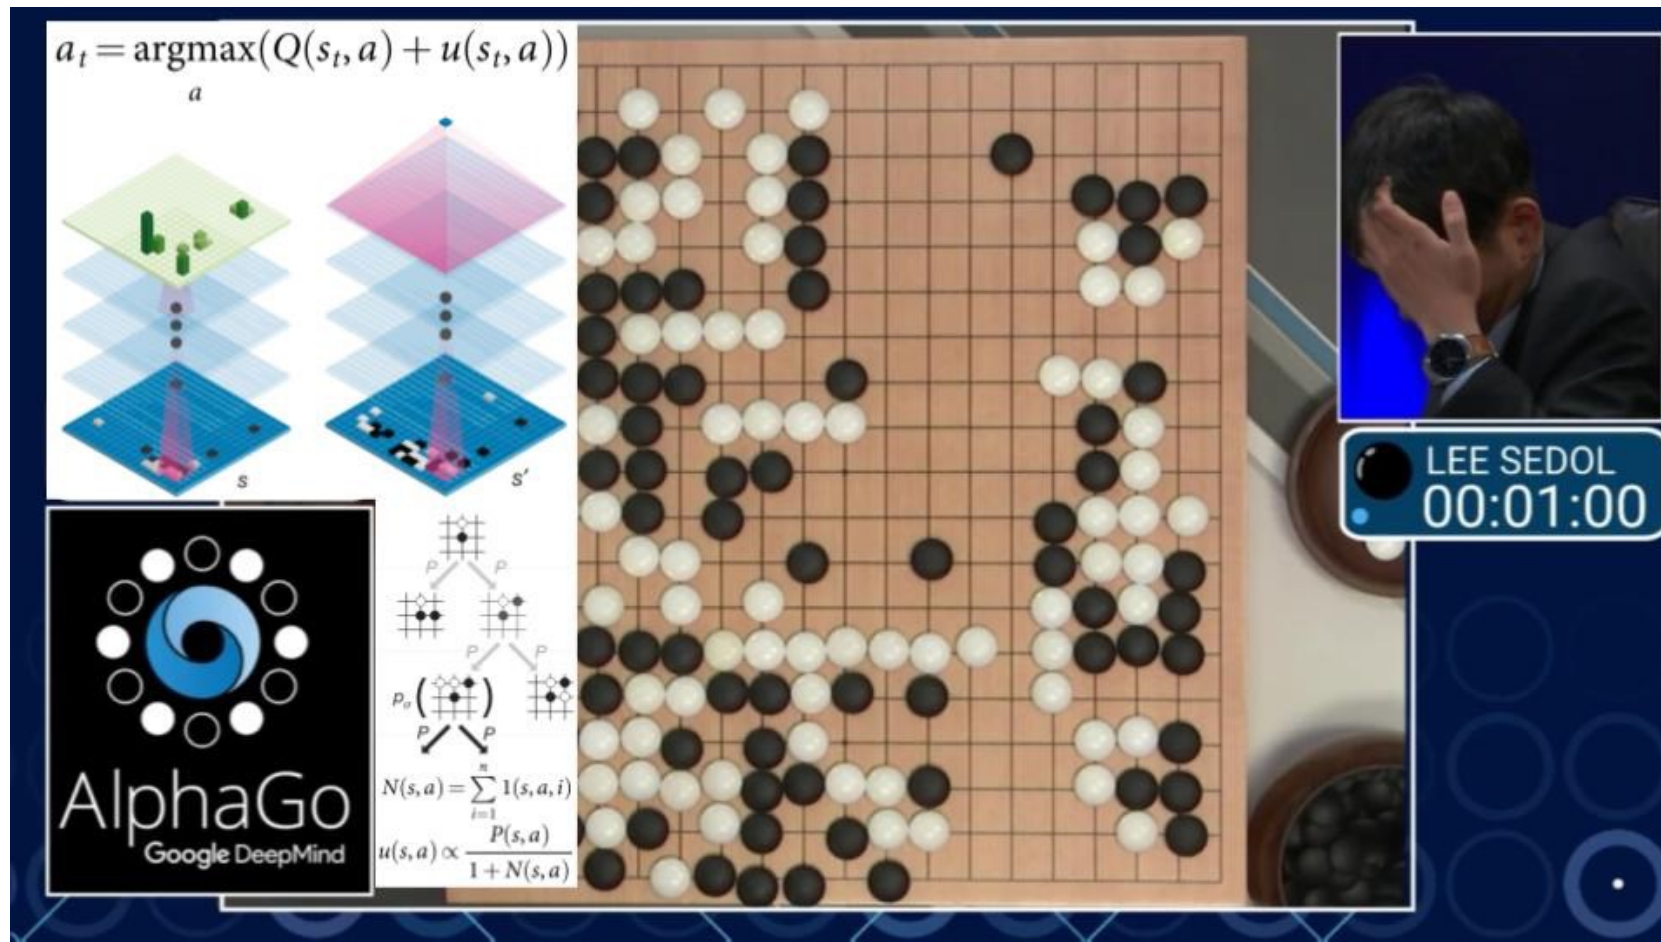


Figure 1: Screen shots from five Atari 2600 Games: (*Left-to-right*) Pong, Breakout, Space Invaders, Seaquest, Beam Rider

... first deep learning model to successfully learn control policies directly from **high-dimensional sensory input** using reinforcement learning.

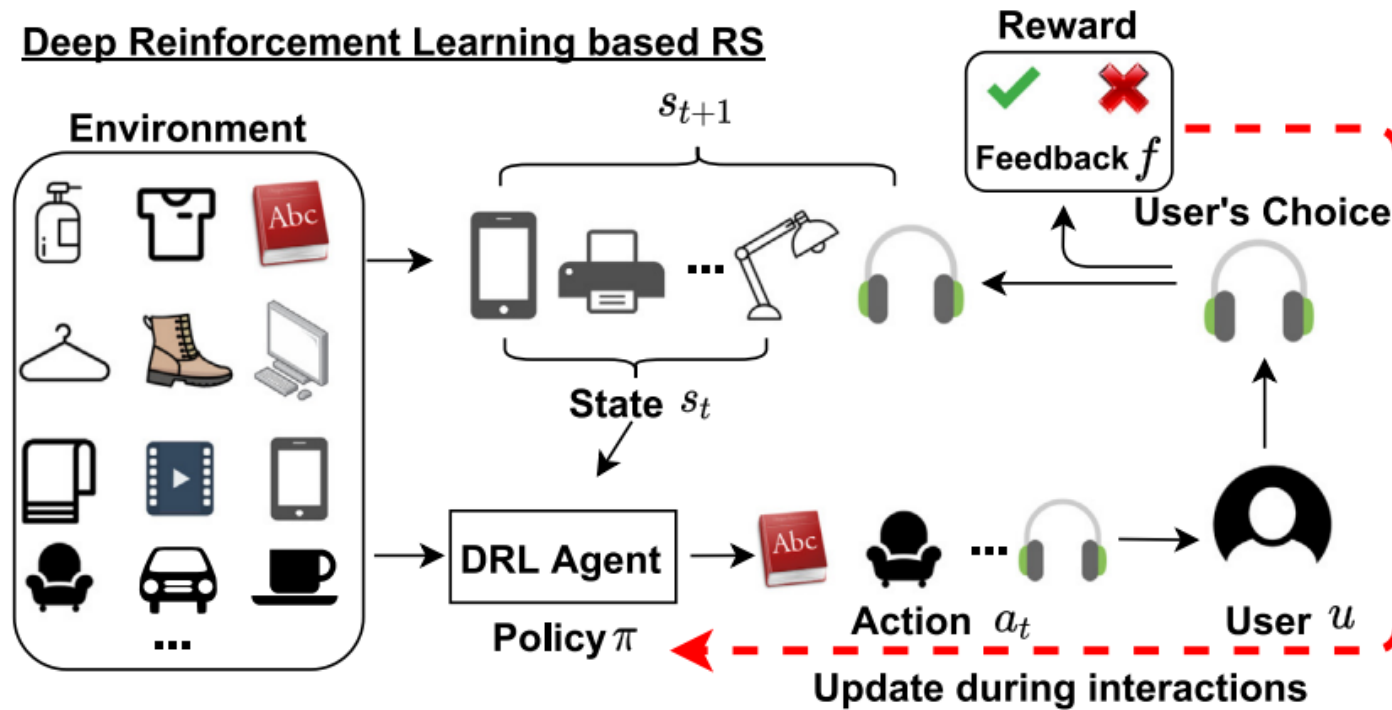
Applications of RL

- David Silver from DeepMind worked on TD algorithm for Go game (2015)



Applications of RL

- Manage investment portfolio
- RL for online learning (suggesting what to show on the stories / Ad)

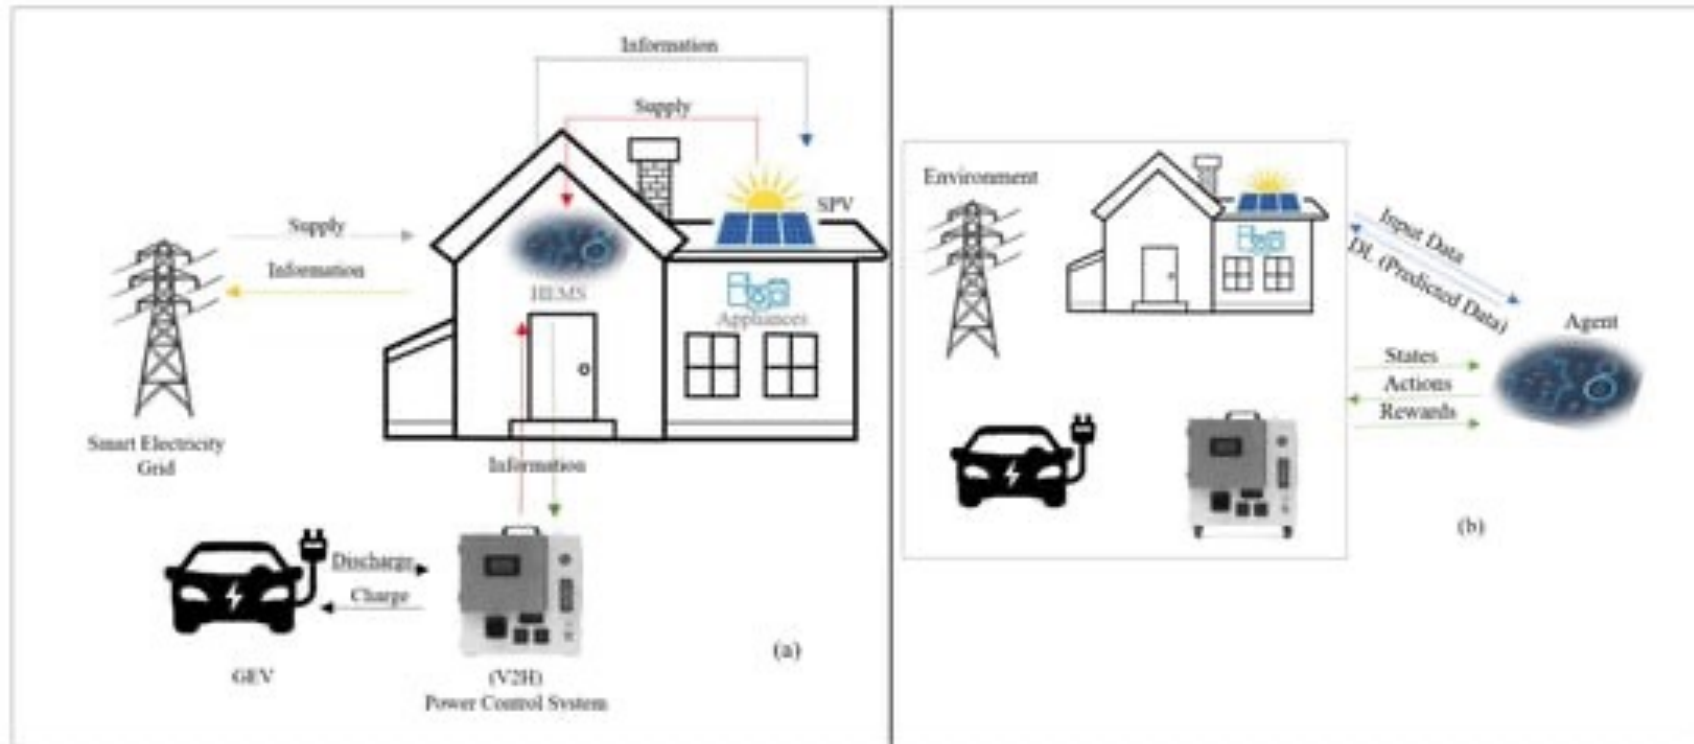


(b) Deep reinforcement learning-based recommender systems

Source: Deep reinforcement learning in recommender systems: A survey and new perspectives by Xiaocong Chen, et al.

Applications of RL

- Control a power station



Source: A Reinforcement Learning Approach for Integrating an Intelligent Home Energy Management System with a Vehicle-to-Home Unit by Ohoud Almugham, et al.

What are Rewards?

- A reward R_t is a scalar feedback signal
- Indicates how well agent is doing at step t
- The agent's goal is to maximize cumulative reward / total future reward

Reward hypothesis in RL:

All goals can be described by the maximization of expected cumulative reward.

- What if there are no intermediate reward?
- What if the task is time sensitive?
- Can the rewards be delayed?
- Is it better to sacrifice immediate reward to gain more long-term reward?

A financial investment may take months to mature.

Refueling a helicopter might prevent a crash in several hours.

What are Rewards?

- Fly stunt maneuvers in a helicopter
 - +ve feedback for following desired trajectory*
 - ve feedback for crashing*
- Defeat the world champion at Backgammon
 - +/-ve feedback for winning/losing a game*
- Manage an investment portfolio
 - +ve feedback for each ₹ in bank*
- Control a power station
 - +ve feedback for producing power*
 - ve feedback for exceeding safety thresholds*
- Make a humanoid robot walk
 - +ve feedback for forward motion*
 - ve feedback for falling over*